

Name – Goutham Selvakumar

Course Name – Neural Networks and Deep Learning

Course Number – CSC 578

Kaggle Username – Goutham Selvakumar

Rank – 10 [246.29115] (Public Score)

Final Class Project (A) Time Series Kaggle

Introduction:

- I created a **Baseline Model** and also **five** other models in order to see which one does better in **performance**. I was able to achieve a better result with the **Model – 4 (Kaggle Best Model)** which will be explained below as we go. So, the other Models were executed and produced results were good and bad at the same time.
- Before conducting the experiments in these models, firstly, I loaded the data and the **preprocessing** steps were done that includes the **Data Cleaning, Data Transformation, and Feature Extraction**.
- Second step, was to use the method **Data Windowing** as from the **Tensor Flow Tutorial**, this was an important task for the **time-series forecasting**. This **Window Generator class** helps to convert the data into a usable format from which the model might be able to learn from the last **12 hours** in order to predict the future data.
- Splitting the Data was done as the next process, it was split into three parts, namely **training, validation, and testing**. Next, I created **dictionaries** in order to store the performance of the model on the **validation and testing** datasets.
- According to the instructions, I created a **Window Generator class** with the respective parameters as the input width – **12**, offset – **3**, and label width – **1**. Using the **Tensor Flow Tutorial**, I was able to create a function that compiles and fits the model and uses the **MSE** as the **loss function** and the **Adam Optimizer** in it.
- An important note is that **I did not shuffle the data before or after the train/test splitting**. So, the next process would be to construct the Baseline Model and the other models in order to see the performance in them.

Models and Development:

BASELINE MODEL

Model – 1:

Architecture: -

- Here, in this baseline model I trained a **LSTM-based model** that has **2 LSTM layers**, **2 dropout layers**, and **3 dense layers**. In order to avoid the **overfitting** in the model I also included the dropout layers on **regularizing** the model.

```
Model: "sequential"
Layer (type)                Output Shape              Param #
-----
lstm (LSTM)                  (None, 12, 128)          78336
dropout (Dropout)            (None, 12, 128)          0
lstm_1 (LSTM)                (None, 12, 128)          131584
dropout_1 (Dropout)          (None, 12, 128)          0
dense (Dense)                (None, 12, 64)           8256
dense_1 (Dense)              (None, 12, 32)           2080
dense_2 (Dense)              (None, 12, 1)            33
reshape (Reshape)            (None, 1, 12)            0
-----
Total params: 220,289
Trainable params: 220,289
Non-trainable params: 0
```

Expectation: -

- From the **epoch**, we can see that the model is learning pretty well with each of the epoch, where the **loss** and **mean absolute error** is getting **reduced** on the **training** and **validation**. This will lead to a **better accuracy** of the model.
- Furthermore, from the **epoch**, the performance on the validation seems to be **irregular**, like we can see there is a small improvement in the **mean absolute error** and **loss** that tells about the overfitting present in it.


Result: -

- In order to further improve the performance of the model and reduce the loss, we are moving onto the other models that helps in improving the performance.
- Among the other models, this model's performance was good with the performance and the value on both the **loss** and **mean absolute error**.


Epoch (50/50):

```
Epoch 50/50
955/955 [=====] - 12s 13ms/step - loss: 0.0147 - mean_absolute_error: 0.0781 - val_loss: 0.0095 - val_mean_absolute_error: 0.0681
156/156 [=====] - 1s 4ms/step - loss: 0.0095 - mean_absolute_error: 0.0681
```

Window Plot:



Loss Plot:



BEST MODEL

Model – 4: Bidirectional RNNs: [Kaggle-Best Model]

- Among all the models that I have experimented, the **loss value** in this **Model -4** seems to be **decreasing** very good consistently which indicates that the model is trained well.

Architecture:-

- This Model – 4 uses **one Bidirectional LSTM model** with **512 units** and **2 LSTM layers** of **128 units**, **2 additional LSTM layers** with **128 units** on each of the layer and then **2 dense layers** that consists of **64** and **32** units respectively. The **output layer** consists of a **1 unit dense layer**.

Model: "sequential_3"

Layer (type)	Output Shape	Param #
bidirectional (Bidirectional LSTM)	(None, 12, 1024)	2199552
lstm_13 (LSTM)	(None, 12, 128)	590336
lstm_14 (LSTM)	(None, 12, 128)	131584
dense_9 (Dense)	(None, 12, 64)	8256
dense_10 (Dense)	(None, 12, 32)	2080
dense_11 (Dense)	(None, 12, 1)	33
reshape_3 (Reshape)	(None, 1, 12)	0

=====
 Total params: 2,931,841
 Trainable params: 2,931,841
 Non-trainable params: 0

Expectation:-

- From the below epochs of **50**, we can conclude that the **loss** on the **training** seems to be **reduced increasingly well** and the **validation loss** is also **reduced** by the number of epochs that indicates that the model is **not overfitting**.
- The **MAE (Mean Absolute Error)** is also getting **reduced** overtime by the number of epochs, since MAE tells how the model's prediction are on average. I trained the LSTM on the sequence data, and the **bidirectional RNNs** here is capable of allowing the past and future data in the sequence to be added.


Result:-

- With this model, I was able to achieve a **significant result** among the **Kaggle [246.29115]** and this was the **Best Model** in comparison with the other models because of **the LSTM Bidirectional LSTM layer**.


Epoch (50/50):

```
Epoch 50/50
955/955 [=====] - 83s 87ms/step - loss: 0.0029 - mean_absolute_error: 0.0337 - val_loss: 0.0042 - val_mean_absolute_error: 0.0415
156/156 [=====] - 4s 25ms/step - loss: 0.0042 - mean_absolute_error: 0.0415
```

Window Plot:



Loss Plot:



WORST MODEL

Model – 2: Adding Layers

Architecture:-

- Due to the addition of **2 LSTM layers** with **128** units each may have been not good as it's contributing to the **model's worst performance**. This model consists of **2 dense layers** with **64** and **32** neurons and the **last dense** having **one** unit respectively with **2 dropout layers** in the model.

Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm_2 (LSTM)	(None, 12, 240)	254400
lstm_3 (LSTM)	(None, 12, 240)	461760
dropout_2 (Dropout)	(None, 12, 240)	0
lstm_4 (LSTM)	(None, 12, 128)	188928
lstm_5 (LSTM)	(None, 12, 128)	131584
dropout_3 (Dropout)	(None, 12, 128)	0
dense_3 (Dense)	(None, 12, 64)	8256
dense_4 (Dense)	(None, 12, 32)	2080
dense_5 (Dense)	(None, 12, 1)	33
reshape_1 (Reshape)	(None, 1, 12)	0

=====
 Total params: 1,047,041
 Trainable params: 1,047,041
 Non-trainable params: 0
 =====

Expectation:-

- The **training** and **validation loss** was way worse in comparison to the other models that I experimented, within each of the epochs I could see the model performing worst.
- This could be due to the **overfitting** of the model, since there are **additional layers** present in the model that could make the model to fit too closely to the training data, which in term causes the performance to be worse.
- Also, the use of **dropouts** which prevents the overfitting by setting a **fraction rate** of units randomly for each training time. These may lead to overfitting of the model if not performing better.


Result:-

- With the **additional layers** the model performed the **worst** among the others, and also the **performance** was not that good.

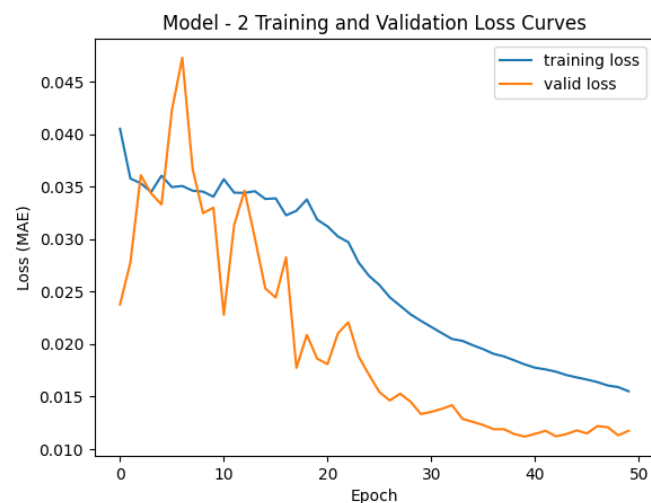
Epochs (50/50):

```
Epoch 50/50
955/955 [=====] - 51s 53ms/step - loss: 0.0155 - mean_absolute_error: 0.0809 - val_loss: 0.0117 - val_mean_absolute_error: 0.0779
156/156 [=====] - 2s 11ms/step - loss: 0.0117 - mean_absolute_error: 0.0779
```

Window Plot:



Loss Plot:



Model – 3: Batch Normalization

Architecture:-

- This model consists of **one LSTM layer** with **240** units and **2 additional LSTM layers** with, **128** units of respectively. Also, **3 BatchNormalization layers**, and **3 dropout layers**. This model also uses an additional of **3 dense layers** with the neurons of **64** and **32** respectively with the last dense layer with **one** unit.

Model: "sequential_2"		
Layer (type)	Output Shape	Param #
lstm_6 (LSTM)	(None, 12, 240)	254400
lstm_7 (LSTM)	(None, 12, 240)	461760
batch_normalization (Batch Normalization)	(None, 12, 240)	960
dropout_4 (Dropout)	(None, 12, 240)	0
lstm_8 (LSTM)	(None, 12, 128)	188928
lstm_9 (LSTM)	(None, 12, 128)	131584
batch_normalization_1 (Batch Normalization)	(None, 12, 128)	512
dropout_5 (Dropout)	(None, 12, 128)	0
lstm_10 (LSTM)	(None, 12, 128)	131584
lstm_11 (LSTM)	(None, 12, 128)	131584
batch_normalization_2 (Batch Normalization)	(None, 12, 128)	512
dropout_6 (Dropout)	(None, 12, 128)	0
dense_6 (Dense)	(None, 12, 64)	8256
dense_7 (Dense)	(None, 12, 32)	2880
dense_8 (Dense)	(None, 12, 1)	33
reshape_2 (Reshape)	(None, 1, 12)	0
Total params: 1,312,193		
Trainable params: 1,311,201		
Non-trainable params: 992		

Expectation:-

- With the end of the **epochs 50/50** we can see that the **loss** and **mean absolute error** was **0.0157** and **0.0748** on the **validation data** which explains that the model is learning well and is performing well after the end of epoch.
- I thought the model would **accurately** run the **training** and predicts the **target variable**, and as we know the model's accuracy improves when there is a **decrease** in the **loss**.


Results:-

- This model performed well in comparison to the other models but not as good as the **Model – 4** that was the **best model**. Because of the **additional LSTM layers**, and the **batch normalization**, the model was able to perform better.


Epochs (50/50):

```
Epoch 50/50
955/955 [=====] - 60s 63ms/step - loss: 0.0050 - mean_absolute_error: 0.0450 - val_loss: 0.0157 - val_mean_absolute_error: 0.0748
156/156 [=====] - 2s 13ms/step - loss: 0.0157 - mean_absolute_error: 0.0748
```

Window Plot:



Loss Plot:



Reaction and Reflection on the results and competition:

- This is the second time I'm joining a kaggle competition. It was challenging and enjoyable at the same time, watching many of my classmates competing each other for a better score. I was able to bring up a good score with my model, but the hardest part was to experiment with different hyperparameters and try to achieve a better score with everyone among the leaderboard was the hardest. My previous scores and the ranking motivated me to bring up with different models and get a better score every time I submit my predictions.

Reaction and Reflection on the project and course:

- With this advanced course, I was able to learn more in-depth about the neural networks, deep learning and with various assignments and quiz I was able to learn more methods in python and also the tensorflow. This is my second advanced course, and with every assignment the difficulty was increasing, overall the course was very challenging since it involved new techniques that I have not implemented and tried to learn it from the beginning and getting the results were challenging.